



UNAM
POSGRADO

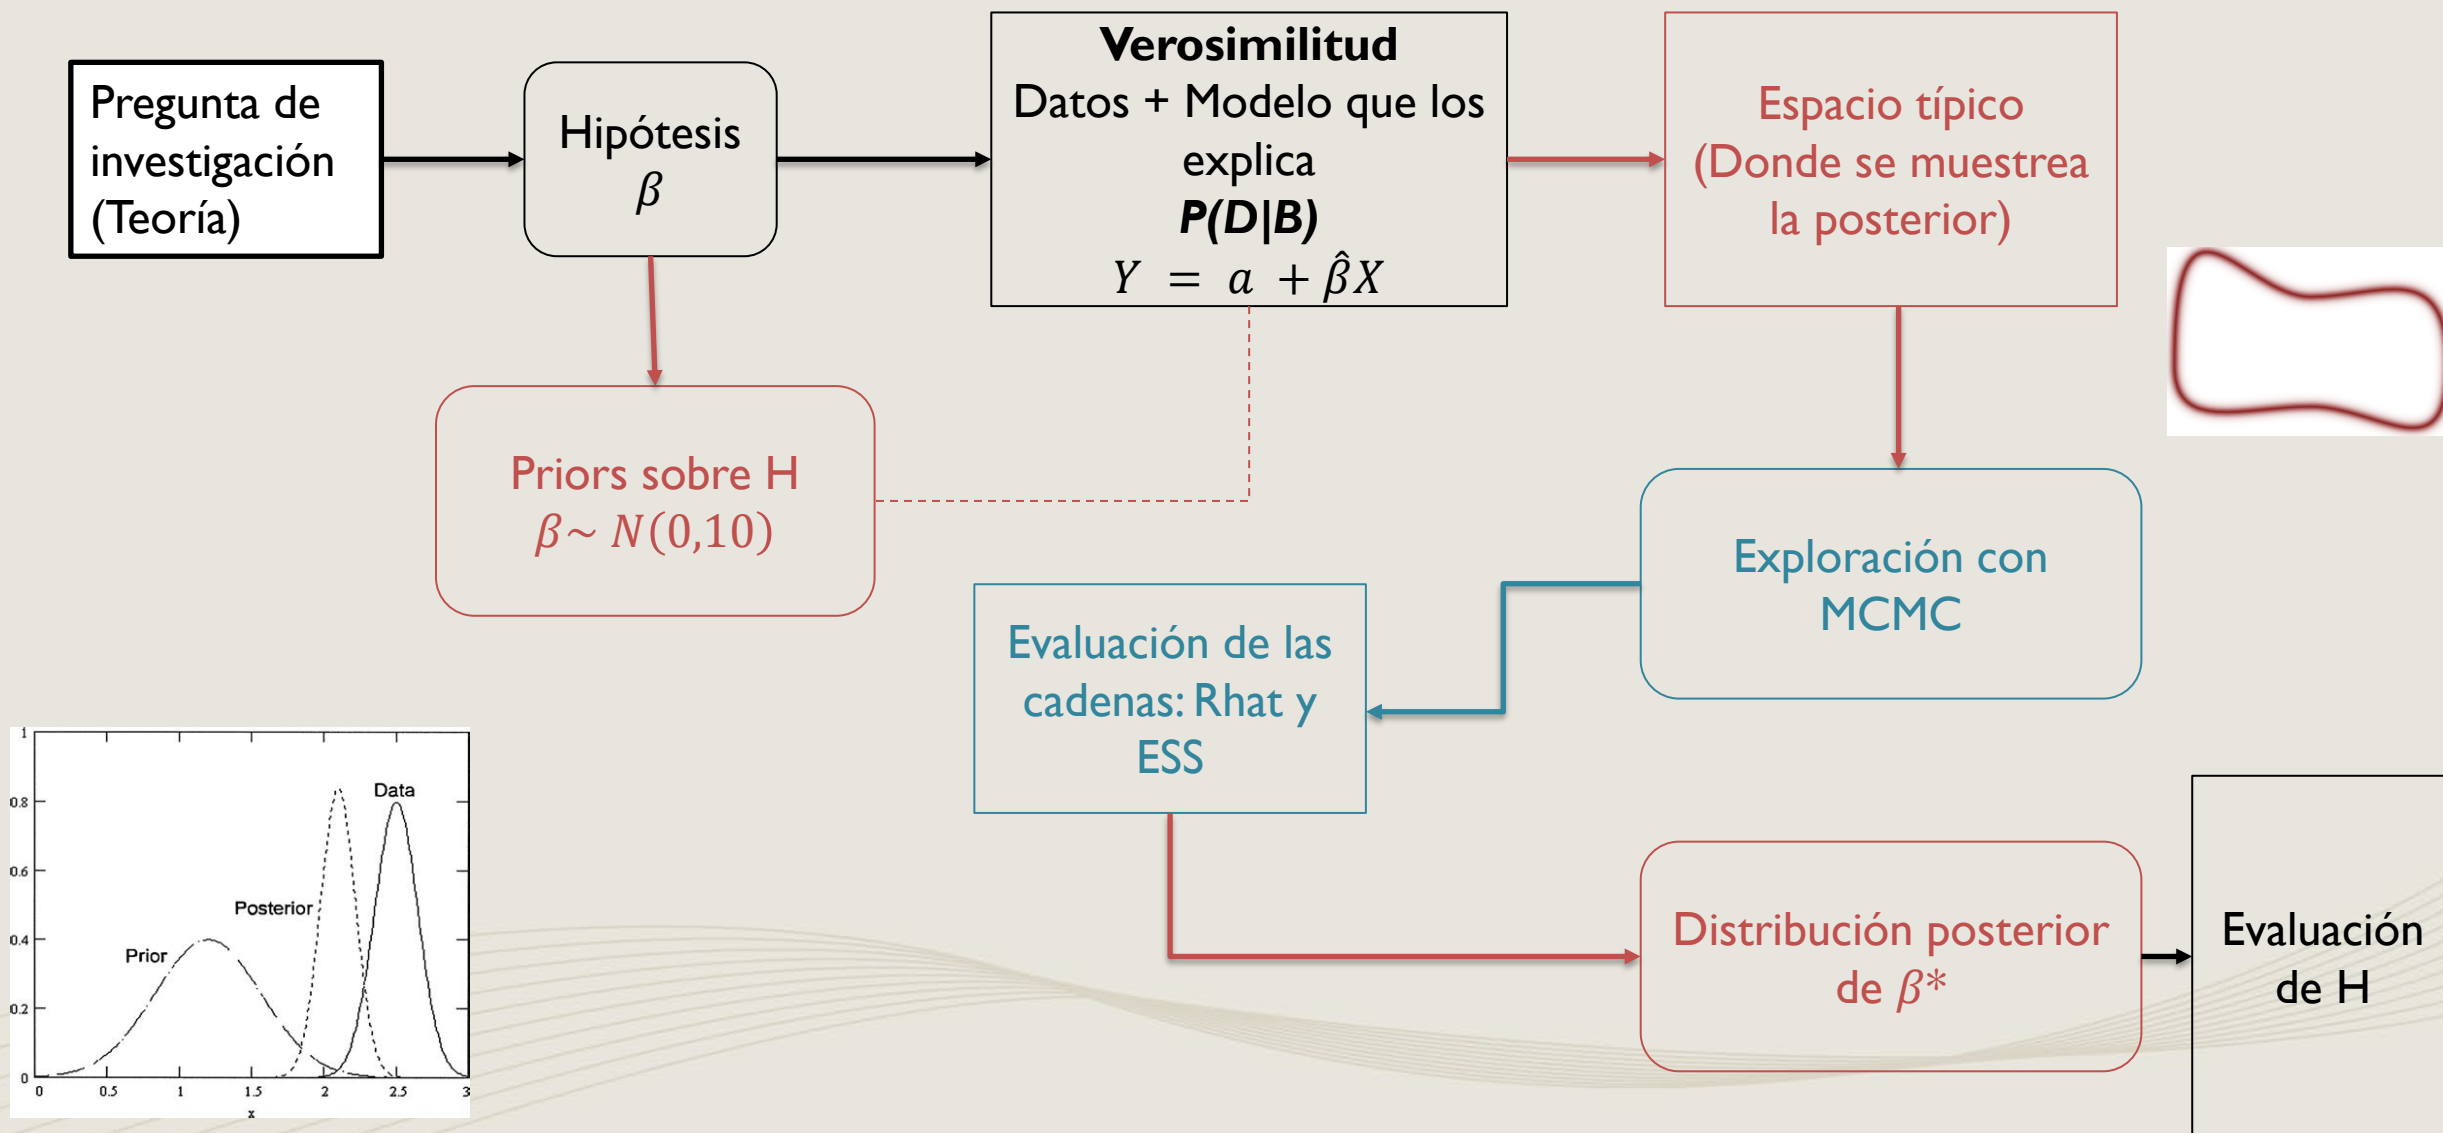


Programa
Universitario
de Estudios
del Desarrollo
UNAM

Estimación en R-Stan

Dr. Héctor Nájera
Dr. Curtis Huffman

Pasos en inferencia bayesiana



Stan: Diferentes plataformas



RStan

The R interface to Stan

Stan está escrito
en C++

Programas

rstan

PyStan

Stan.jl

CmdStan

Rstan (crudo)

rstan
(interfaces
para rstan)

rstanarm

rethinking

brms

cmdstanr

Tidybayes

bayesplot

Avanzado: mayor potencia y flexibilidad

Intermedio: menor flexibilidad

Intermediarios: brms y rstanarm

Además brms conecta con cmdstanr y puede ser multicore y multithread –y ahora se puede usar con las tarjetas gráficas- dentro de cada cadena!

Noten las combinaciones posibles:

- Modelo espacio/temporal (AR y CAR) inflado de ceros!!!!

brms offers more modeling capabilities, flexibility with priors, and more³

| | brms | rstanarm |
|-----------------------------------|------------------|-------------------------|
| Supported model types: | | |
| Linear models | yes | yes |
| Robust linear models | yes | yes ¹ |
| Binomial models | yes | yes |
| Categorical models | yes | no |
| Multinomial models | no | no |
| Count data models | yes | yes |
| Survival models | yes ² | yes |
| Ordinal models | various | cumulative ³ |
| Zero-inflated and hurdle models | yes | no |
| Generalized additive models | yes | yes |
| Non-linear models | yes | no |
| Additional modeling options: | | |
| Variable link functions | various | various |
| Weights | yes | yes |
| Offset | yes | yes |
| Multivariate responses | limited | no |
| Autocorrelation effects | yes | no |
| Category specific effects | yes | no |
| Standard errors for meta-analysis | yes | no |
| Censored data | yes | no |
| Truncated data | yes | no |
| Customized covariances | yes | no |
| Missing value imputation | no | no |
| Bayesian specifics: | | |
| parallelization | yes | yes |
| population-level priors | flexible | normal, Student-t |
| group-level priors | normal | normal |
| covariance priors | flexible | restricted ⁵ |
| Other: | | |
| Estimator | HMC, NUTS | HMC, NUTS |
| Information criterion | WAIC, LOO | AIC, LOO |
| C++ compiler required | yes | no |
| Modularized | no | no |

Spatial models

Si no se puede hay que usar rstan

¿Intermediario?

- Brms:

```
MI.I<-brm(ins_ali_mo_3 ~ ed_jefe_2 + ed_jefe_3 + ed_jefe_4 + ed_jefe_5
+ sexo_jefe_1 + edad_jefe + niv_ed_prom + tv + compu +
refri + lavad + tvpaga + hablaind + tinaco + internet + tot_integ +
salario_mun +
i_nolav + i_nostvp + i_noref + salario_ent + rural + (1|id_mun) +
(1|clave_ent),
data=DI, family = bernoulli(),
prior = c(set_prior("normal(0,1)", class = "b"),
set_prior("normal(0,5)", class= "sd")),
warmup = 400, iter = 1800, control = list(adapt_delta = .96),
chains = 2, cores=14, threads = threading(14, grainsize = 100),backend =
"cmdstanr")
```

STAN:

```
functions {
  /* integer sequence of values
  * Args:
  * start: starting integer
  * end: ending integer
  * Returns:
  * an integer sequence from start to end
  */
  int[] sequence(int start, int end) {
    int seq[end - start + 1];
    for (n in 1:num_elements(seq)) {
      seq[n] = n + start - 1;
    }
    return seq;
  }
  // compute partial sums of the log-likelihood
  real partial_log_lik_lpmf(int[] seq, int start, int end, data int[] Y, data matrix Xc, vector b, real Intercept, data int[] J_1, data vector
  Z_1_1, vector r_1_1, data int[] J_2, data vector Z_2_1, vector r_2_1) {
    real ptarget = 0;
    int N = end - start + 1;
    // initialize linear predictor term
    vector[N] mu = Intercept + rep_vector(0.0, N);
    for (n in 1:N) {
      // add more terms to the linear predictor
      int nn = n + start - 1;
      mu[nn] += r_1_1[J_1[nn]] * Z_1_1[nn] + r_2_1[J_2[nn]] * Z_2_1[nn];
    }
    ptarget += bernoulli_logit_glm_lpmf(Y[start:end] | Xc[start:end], mu, b);
    return ptarget;
  }
}
data {
  int<lower=1> N; // total number of observations
  int Y[N]; // response variable
  int<lower=1> K; // number of population-level effects
  matrix[N, K] X; // population-level design matrix
  int grainsize; // grainsize for threading
  // data for group-level effects of ID 1
  int<lower=1> N_1; // number of grouping levels
  int<lower=1> M_1; // number of coefficients per level
  int<lower=1> J_1[N]; // grouping indicator per observation
```



STAN Cont.

```
• STAN
• // group-level predictor values
• vector[N] Z_1_1;
• // data for group-level effects of ID 2
• int<lower=1> N_2; // number of grouping levels
• int<lower=1> M_2; // number of coefficients per level
• int<lower=1> J_2[N]; // grouping indicator per observation
• // group-level predictor values
• vector[N] Z_2_1;
• int prior_only; // should the likelihood be ignored?
• }
• transformed data {
•   int Kc = K - 1;
•   matrix[N, Kc] Xc; // centered version of X without an intercept
•   vector[Kc] means_X; // column means of X before centering
•   int seq[N] = sequence(1, N);
•   for (i in 2:K) {
•     means_X[i - 1] = mean(X[, i]);
•     Xc[, i - 1] = X[, i] - means_X[i - 1];
•   }
• }
• parameters {
•   vector[Kc] b; // population-level effects
•   real Intercept; // temporary intercept for centered predictors
•   vector<lower=0>[M_1] sd_1; // group-level standard deviations
•   vector[N_1] z_1[M_1]; // standardized group-level effects
•   vector<lower=0>[M_2] sd_2; // group-level standard deviations
•   vector[N_2] z_2[M_2]; // standardized group-level effects
• }
• transformed parameters {
•   vector[N_1] r_1_1; // actual group-level effects
•   vector[N_2] r_2_1; // actual group-level effects
•   r_1_1 = (sd_1[1] * (z_1[1]));
•   r_2_1 = (sd_2[1] * (z_2[1]));
• }
• model {
•   // likelihood including constants
•   if (!prior_only) {
•     target += reduce_sum(partial_log_lik_lpmf, seq, grainsize, Y, Xc, b, Intercept, J_1, Z_1_1, r_1_1, J_2, Z_2_1, r_2_1);
•   }
•   // priors including constants
•   target += normal_lpdf(b | 0, 1);
•   target += student_t_lpdf(Intercept | 3, 0, 2.5);
•   target += normal_lpdf(sd_1 | 0, 5)
```


¿Por dónde empezar brms o stan?

La escritura de los modelos bayesianos tiene la ventaja de que es necesario saber lo que uno está haciendo

Escribir el modelo estadístico no es necesariamente saber el modelo estadístico

Los LLM (DeepSeek, ChatGpt, ...) facilitan la escritura/traducción del código pero eso es insuficiente para hacer inferencias válidas

El paquete brms es genial porque es una interface muy amigable para estimar modelos bayesianos en stan pero debe enseñarse después y no antes de la escritura en stan

En el mercado de trabajo va a importar menos que puedan usar brms y más que sepan stan.

La gente más buscada para estimar modelos es la que tenga acceso al mayor poder en LLMs y el conocimiento más sólido de stan



Estructura de modelo bayesiano

Estructura de un modelo Bayesiano:

$\text{outcome}_i \sim \text{Normal}(\mu_i, \sigma)$
 $\mu_i = \beta \times \text{predictor}_i$
 $\beta \sim \text{Normal}(0, 10)$
 $\sigma \sim \text{HalfCauchy}(0, 1)$

$D_i \sim \text{Normal}(\mu_i, \sigma)$
 $\mu_i = \alpha + \beta_A A_i$
 $\alpha \sim \text{Normal}(10, 10)$
 $\beta_A \sim \text{Normal}(0, 1)$
 $\sigma \sim \text{Uniform}(0, 10)$

| | |
|--|------------------------|
| $D_i \sim \text{Normal}(\mu_i, \sigma)$ | [likelihood] |
| $\mu_i = \alpha + \beta_R R_i + \beta_A A_i$ | [linear model] |
| $\alpha \sim \text{Normal}(10, 10)$ | [prior for α] |
| $\beta_R \sim \text{Normal}(0, 1)$ | [prior for β_R] |
| $\beta_A \sim \text{Normal}(0, 1)$ | [prior for β_A] |
| $\sigma \sim \text{Uniform}(0, 10)$ | [prior for σ] |

<https://mc-stan.org/docs/stan-users-guide/linear-regression.html>

Pasos para la modelación



RStan

The R interface to Stan

- Programa: R + Rstudio
- Algoritmo HMC: Rstan
- Interface: brms + cmdstan



CmdStanR



Paso I

- Instalar (asegurarse) R version 4 o superior:

<https://cran.r-project.org/bin/windows/base/>



Paso 2

- Instalar Rstudio

<https://www.rstudio.com/products/rstudio/download/#download>



Paso 3

- Instalar Rtools 4.0 o superior

<https://cran.r-project.org/bin/windows/Rtools/>



Paso 4

- Instalar Rstan

Seguir estos pasos

<https://github.com/stan-dev/rstan/wiki/RStan-Getting-Started>

Asegurarse que el ejemplo “corre”



Paso 5

- Instalar en Rstudio

Brms, tidybayes, bayesplot, posterior, tidyverse



- Instalar cmdstanr –mayor rapidez–

<https://mc-stan.org/cmdstanr/>



Introducción –refrescando- R Software

- Algunas guías
 - https://www.youtube.com/watch?v=_V8eKsto3Ug



GitHub

- <https://github.com/hectornajera83/ClaseBayes2025>